

000
001
002
003
004
005
006
007
008
009
010
011054
055
056
057
058
059
060
061
062
063
064
065
066
067
068
069
070
071
072
073
074
075
076
077
078
079
080
081
082
083
084
085
086
087
088
089
090
091
092
093
094
095
096
097
098
099
100
101
102
103
104
105
106
107

Fast-Match: Fast and robust feature matching on large images

Anonymous CVPR submission

Paper ID ****

Abstract

Both consumer cameras and camera phones produce images that often exceed 10 megapixels. Yet computing and matching local features in images of this size can easily take more than twenty seconds using fast matching algorithms. This is much too slow for interactive applications and much too expensive for large scale image operations. We introduce Fast-Match, an algorithm designed to swiftly match large images without compromising on matching precision or recall. Fast-Match derives its speed from only computing features in those parts of the image that can be confidently matched. We show that Fast-Match is an order of magnitude faster than Ratio-Match, yet often doubles the precision on difficult to match image pairs at equal recall. In addition we prove that when one image is known in advance, Fast-Match can achieve linear performance $O(n)$ in the number of feature points n .

1. Introduction

Whenever we match local image features we are faced with a choice between performance and precision. On one hand SIFT features proposed by Lowe [22] have shown again and again to compare favorably to other local image descriptors, especially under unconstrained conditions such as perspective change with non-planar scenes [15, 23, 24]. On the other, SIFT keypoints and descriptors are slow to compute, the main raison d'être for the introduction of various alternative local image features. In many applications of computer vision we would like the increase the computational performance in order to work on larger images, bigger image sets, at faster frame rates or with more limited hardware, but we cannot give up the additional precision that SIFT affords us over other local features without negative effects for our application.

In this paper we introduce a matching algorithm designed to match features between two images only in image areas that are likely to correspond. This approach is much faster than traditional methods because there is no need to compute descriptors for areas in the image that are

not matched. We provide two variations of the algorithm: The *general* variation functions as a traditional matching algorithm and matches two unknown images albeit faster and more robustly than conventional methods. The *retrieval* variation on the other hand assumes that we know one of the images we intend to match beforehand; under this assumption, it can be a magnitude faster than existing matching methods without compromising on accuracy. We will unimaginatively refer to the proposed algorithm as *Fast-Match* in this paper and make clear from the context whether this refers to the *retrieval* or *general* variation.

The problem that *Fast-Match* attempts to solve is two-fold. By matching only image areas that are likely to correspond we hope to improve accuracy by entirely ignoring parts of the images that would otherwise likely be a source of incorrect correspondences. At the same time, this enables us to improve computational speed by not having to compute keypoints and descriptors for large parts of the images and at the same time reducing the amount of feature points we need to match. Both problems have been addressed in part by past work.

Fast-Match makes use of an angular assumption to efficiently find new matches in the geometric neighbourhood of already confirmed matches. However this constraint is only applied *locally* to increase the number of matches, making *Fast-Match* robust to outliers. In addition *Fast-Match* does not rely on an initial set of matches and derives much of its speed from the fact that not even an initial set of local image features is required.

This paper is structured as follows: Section 2 discuss related work. Section 3 introduces the *Fast-Match* algorithm. Section 4 outlines the experimental setup. In Section 5 we discuss the results before concluding in Section 6.

2. Related work

As noted, the cost of computing SIFT keypoints and descriptors is addressed by other local image features such as SURF [3], BRIEF [6], or BRISK [19], just to mention a few. Similarly efforts have been made to improve SIFT itself such as PCA-SIFT [16] and GLOH [23]. Both apply PCA to reduce descriptor lengths and improve distinctive-

108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
ness, but neither have been shown to consistently outperform SIFT [20, 23].

132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
Efforts to reduce the computational costs of finding nearest neighbours to feature points have largely focused on metric trees. Naively the set of nearest neighbours between features in two images can be computed by brute force in $O(n^2)$, where n is the total number of feature points in the two images (a typical three megapixel image may contain anywhere from 500 to 5000 feature points). When SIFT was originally published Lowe, proposed using the Best-Bin-First method to approximately search for nearest neighbours [4, 21]. This reduces the computational complexity to $O(n \log n)$, but even approximate metric trees are hard pressed to compete with brute force due to the high dimensionality of SIFT and the constant costs incurred with constructing and searching in a metric tree. Later work by Muja and Lowe [25] focused on improving approximate nearest neighbour searches by using several KD-Trees simultaneously while optimizing the tree structure using K-Means to cluster similar features. Recent work on knn-graphs shows a lot of promise for high dimensional cases [13]. We later review these improvements and their effect on efficiently matching large scale images.

154
155
156
157
158
159
160
161
A wealth of matching methods have focused instead on increasing the efficiency of local feature matching. *Fast-Match* builds upon the foundation of *Ratio-Match* introduced originally by Deriche et al. [12] and Baumberg [2] even though Lowe is usually credited for introducing it [22]. They both propose using the ratio of the similarity of the best to second best match of a given point to evaluate how unique it is. Their finding has later been tested by several independent teams, all concluding that thresholding based on this ratio is generally superior to thresholding based on similarity or returning all nearest neighbours [22–24, 28].

162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
Brown and Lowe [5] extend ratio match to deal with a set of images by using not the ratio of the best and second best match, but the average ratio of the best and the average of second best matches across a set of images. Rabin et al. [28] try to enhance descriptor matching by looking at the statistical distribution of local features in the matched images, and only return a match when such a correspondence would not occur by mere chance. Finally, Arnfred et al. generalize *Ratio-Match* to make use of both of the matched images to provide a more accurate evaluation of the uniqueness of a match [1].

For sparse local image features many solutions have combined *Ratio-Match* with various geometric constraints to improve matching. These constraints are based on assumptions regarding the transformation between the *query* and *target images*. A commonly used example is *RANSAC* applied to feature matching where matches are chosen from a pool of candidates according to how well they approximate a global epipolar geometry [14, 30, 31]. Similar global



Figure 1: The 50 best matches by the proposed *Fast-Match* algorithm (left) and the *Ratio-Match* algorithm [22] (right). The lines between the image pairs denote a match as proposed by the algorithm. For *Fast-Match* the zones of the right image where local features were computed has been marked with a blue square.

angular and distance constraints can be used to filter a set of matches as shown in [17, 29]. Finally the problem of feature matching can be modelled as a graph matching problem where each feature is a vertex, and edge values correspond to a geometric relation between two features as demonstrated in [9, 11, 27, 32, 34]. For cases that require more flexibility such as scenes with moving elements and non-planar perspective change Pairwise constraints provide an alternative to the more rigid global constraints. Introduced by Leordeanu and Herbert [18], pairwise constraints work by applying a geometric constraint across matches on a pairwise basis, attempting to minimize the pairwise error. This approach has later been adopted by Pang et al. but made scale invariant and more efficient by using a voting scheme [26, 35]. Similarly we can cluster matches together based on a geometric constraint and treat each cluster as a separate case [8, 10, 33]. An interesting application of this principle is demonstrated by Chen et al. who iteratively use a Hough transform to cluster matches using angular constraints [7].

While geometric constraints have been shown to work well, they are often susceptible to outliers and tend to be computationally demanding. All of the above geometric methods require a set of initial matches usually provided by *Ratio-Match* which acts as a lower bound on their running time. In practice even fast graph matching methods such as the covering trees method proposed by Yarkony et al. are two or three magnitudes slower than *Ratio-Match* [34].

3. Matching Fast and Slow

In this section we will introduce the fundamentals of *Fast-Match* and motivate the design choices as we go. The algorithm consists of 3 components; finding seeds, finding matches, and exploring for other places where matches might be.

3.1. Considerations

If we set out to design a truly fast matching algorithm, we cannot rely just on optimizing the matching step once

216 the descriptors from both images have been found and computed.
 217 Finding and computing descriptors can easily account for 80% of the time spent matching for bigger images
 218 (cf. Figure 7). For this reason *Fast-Match* is designed to
 219 only compute features for the part of the image we hope to
 220 match.
 221

222 The *Fast-Match* algorithm is demonstrated in Algo-
 223 rithm 1. Given a *query image* and a *target image* that we
 224 intend to match and a confidence threshold τ , we obtain a
 225 set of seed matches from the two images. For each seed
 226 match we look at the matched position in the *query* and *tar-
 227 get images* and find a set of matches. We save these matches
 228 and their confidence scores; for those that pass the con-
 229 fidence threshold τ , we obtain another set of seed matches.
 230 In this way we iterate until we have no more seed matches
 231 and return the matches and their confidence scores. In an in-
 232 tuitive sense τ serves as a parameter directing the thorough-
 233 ness of the algorithm, i.e. how much time we spend match-
 234 ing, while the final precision and recall can be adjusted by
 235 thresholding the matches on their confidence scores at the
 236 end.
 237

238 We a *general* and a *retrieval* variation of the algorithm.
 239 The latter assumes that we know one of the images that
 240 we intend to match ahead of time and can do some off-
 241 line computations. For the *general* variation we make those
 242 computations on the fly instead and do not assume anything
 243 to be pre-computed.
 244

Algorithm 1 Fast-Match

245 **Require:** $I_{\text{query}}, I_{\text{target}}$: images, $\text{maxiter} \in \mathbb{N}$, $\tau \in [0, 1]$
 246 $M_{\text{seed}} \leftarrow \text{seed_matches}(I_{\text{query}}, I_{\text{target}})$
 247 $M_{\text{final}} \leftarrow \emptyset$
 248 $C_{\text{final}} \leftarrow \emptyset$
 249 $M_{\text{seen}} \leftarrow \emptyset$
 250 **while** $M_{\text{seed}} \neq \emptyset \wedge i < \text{maxiter}$ **do**
 251 $M_{\text{round}} \leftarrow \text{get_matches}(M_{\text{seed}})$
 252 $C_{\text{round}} \leftarrow \text{get_confidence}(M_{\text{round}})$
 253 $M_{\text{seed}} \leftarrow \text{get_seeds}(M_{\text{round}} \setminus M_{\text{seen}}, C_{\text{round}}, \tau)$
 254 $M_{\text{seen}} \leftarrow M_{\text{seen}} \cup M_{\text{seed}}$
 255 $M_{\text{final}} \leftarrow M_{\text{final}} \cup M_{\text{round}}$
 256 $C_{\text{final}} \leftarrow C_{\text{final}} \cup C_{\text{round}}$
 257 **end while**
 258 **return** $M_{\text{final}}, C_{\text{final}}$

3.2. Initiating Seeds

263 Depending on the scenario, several strategies can be used
 264 to obtain a set of initial seed matches. In practice we have
 265 found it efficient to resize both images to thumbnails and
 266 use *Ratio-Match* to obtain a set of matches and ratios that
 267 we then threshold with the confidence threshold τ to obtain
 268 the initial seed matches. The thumbnails created need to be
 269 sized so they on one hand ensure that objects of interest are

270 still detectable by a keypoint detector while on the other re-
 271 main small to ensure that we find matches efficiently. We
 272 determine empirically that a thumbnail size of 300×300
 273 pixels maintains a good balance between speed and accu-
 274 racy. We found this size to be independent of the original
 275 image size as long as we are matching images that are big-
 276 ger than one megapixel.
 277

3.3. Collecting Matches and Computing Confidence

278 If a seed match yields a connection between two points
 279 p_q in the *query image* and p_t in the *target image*, then we are
 280 interested in collecting all matches between the regions R_q
 281 and R_t centered around p_q and p_t respectively. From each
 282 region we can extract a set of feature points between which
 283 we try to find a set of matches M_{qt} and a set of confidence
 284 scores C_{qt} .
 285

286 Lowe and others have shown that the distance between
 287 two SIFT descriptors is much less indicative of a true corre-
 288 spondence than the ratio between the best and second best
 289 match [22–24, 28]. This ratio is more formally defined as
 290 follows: If we let f_q be a feature in the *query image* and
 291 f_t, f_b be the two nearest neighbours of f_q in the *target im-
 292 age* then the ratio r is defined as follows:
 293

$$r = \frac{d(f_q, f_t)}{d(f_q, f_b)}.$$

294 Here $d(f_i, f_j)$ is the distance between the features f_i and f_j .
 295 For SIFT this is the Euclidean distance. The distance be-
 296 tween two feature points is greater when the feature points
 297 resemble each other less. For this reason we have higher
 298 confidence in a match with a low r -value. Using r as a
 299 measure of match confidence presumes that we expect each
 300 feature in the *target image* to have at most one true cor-
 301 respondence in the *target image*. Intuitively if we try to
 302 find a match for a feature f_i in an image that does not have
 303 any true correspondences, then we would expect the two
 304 closest neighbours to be roughly equally well matched with
 305 f_i . The effectiveness of *Ratio-Match* over a naïve nearest
 306 neighbor match ([22, 23]) supports this notion. For this
 307 reason we attribute high confidence to matches where the
 308 closest neighbour is dramatically closer to f_i than the sec-
 309 ond closest neighbour and discard the rest.
 310

311 We are faced with a problem when applying this tech-
 312 nique to obtain the set of confidence scores C_{qt} , since for
 313 any match in M_{qt} we only know the nearest neighbours
 314 amongst the features of R_q and R_t . To get around this
 315 problem we either assume to know one of the images be-
 316 forehand (the *retrieval* variation) or – in case we cannot
 317 make that assumption – we compute the features of one of
 318 the images (the *general* variation). For a given match be-
 319 tween features f_q and f_t we can now find the second clos-
 320 est neighbour f_b and calculate the confidence as
 321 $r = d(f_q, f_t)/d(f_q, f_b)$.
 322

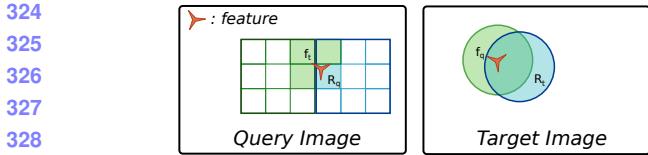


Figure 2: Exploration of features based on match. The areas shaded blue are the areas that were searched to obtain the match. The areas shaded green are candidate areas for obtaining more matches.

In the retrieval scenario where we know an image beforehand, we can further optimize this step. If we assume the *target image* is known in advance we can approximate the ratio by letting $\hat{r} = d(f_t, f_q)/d(f_t, f_b)$. Here the feature f_b is also part of the target image, which means we can pre-calculate $d(f_t, f_b)$ for all features in the cached *target image* before we start matching. When these distances are tied to their corresponding feature in the target image, it becomes trivial to calculate \hat{r} .

3.4. Exploring for Matches

In each iteration we compute a new set of seed matches, i.e. positions that might yield more matches in the image. For each region R_i evaluated during the collection step we now have a set of matches M_i and a set of confidence levels C_i that we can use to predict whether the neighbourhood of R_i is worth exploring.

There are many possible heuristics for predicting possible regions including local and global epipolar assumptions and partial graph isomorphisms. However, for the sake of simplicity and speed we have chosen a straightforward approach based on weak angular assumptions, as illustrated in Figure 2. Assuming that the *target image* has been cached either in advance or before the matching step, the blue shaded areas show R_q and R_t for a given seed match. In the *target image* we collect all features in a given radius while we compute features in the rectangular R_q in the *query image*. For performance reasons we compute all features in the blue square but match only the features inside the shaded area. Next a match is found in the collection step between f_q and f_t . Based on the position of f_q in R_q we select three areas with potential for more matches. The center of each is matched with the center position f_t to produce three seed matches for the next iteration. In Figure 1 we illustrate this process in the *Fast-Match* image pair. The *query image* to the right in the figure has a small blue square for each region that has been used while matching the two images.

In practice it is necessary that these squares overlap in order to detect features lying close to the edges. This incurs a bit of overhead which is minimized by only collecting features for groups of 9 squares. In order to avoid computing

the same matches or features twice, quite a bit of care has to be expended making sure that results are properly cached. A matrix containing ‘bins’ of features can conveniently be used to store features from different image regions. Similarly a hash-set is suitable for keeping track of which regions have already been matched and which matches have already been found.

3.5. Computational Complexity

The difference between the *general* and *retrieval* variation of *Fast-Match* consists in whether we compute the features in the *target image* (general) or if we presume the features are computed offline (retrieval). Computing the features is linear in the number of feature points, while finding $d(f_q, f_b)$ for all n features in the *target image* can be done in $O(n \log n)$ using metric trees.

Once the features and their distances have been computed, the algorithm iteratively finds new seed matches based on previous sets of seed matches. For each seed match we collect new matches, calculate confidence scores, and obtain new seed matches. Each of these steps varies only with the local region size which is constant. As a consequence the running time is linear in the amount of possible seed matches.

We will show that the number of possible seed matches is on average linear in the number of image features, and we also provide an upper bound for the probability that it is not. We assume that outside of true correspondences, features have better or at least equally good matches in terms of confidence in the image they come from. More rigorously put: For any feature in the *target image* f_t let f_1, f_2, \dots be the best, second best, etc matching feature from either image which is not a true correspondence. Let A_{ti} be the event that f_i is found in the *query image*, then $\mathbb{P}(A_{ti}) \leq 0.5$ and A_{ti} is independent from A_{tj} when $i \neq j$.

For a given feature in the *target image* f_t , its nearest neighbour in the *target image* f_b , and the set of features in the *query image* F_{query} , we can define the stochastic variable X_t as follows:

$$X_t = |\{f_q \mid f_q \in F_{query}, d(f_q, f_t) < d(f_q, f_b)\}| \quad (1)$$

That is, X_t is the number of features in the query image closer to f_q than f_b . For each feature f_t , X_t is an i.i.d. stochastic variable.

In the worst case, each feature in the *target image* has a true correspondence in the *query image* and $\mathbb{P}(A_{ti} = 0.5)$, in which case we find can find $\mathbb{P}(X_t = n) = \mathbb{P}(A_{ti})^n = 0.5^n$, as well as the expected value $\mathbb{E}(X_t) = \sum_{i=1}^{\infty} i 0.5^i = 2$, and the variance $Var(X_t) = \sum_{i=1}^{\infty} 0.5^i (i - 2)^2 = 6$.

For $\tau = 1$ we accept a seed match when $d(f_t, f_q) \leq d(f_t, f_b)$. That means that for any feature in the *query image*, we accept at most X_t matches. To find the total amount of seed matches given n , we let $S_n = X_1 + \dots + X_n$, in

324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431

432 which case $\mathbb{E}(S_n) = 2n$. This proves that the expected
 433 amount of seed matches is linear in the number of target
 434 features.

435 For the above case it is theoretically possible that for every
 436 one of the n features we end up considering kn seed
 437 matches for some constant k , i.e. a quadratic behaviour. Using
 438 Chebyshev's inequality we can provide an upper bound
 439 on the probability of this event:

$$441 \quad \mathbb{P} \left(\left| \frac{S_n}{n} - \mathbb{E}(X_t) \right| \geq kn \right) \leq \frac{\text{Var} \left(\frac{S_n}{n} \right)}{k^2 n^2} = \frac{6}{k^2 n^3} \quad (2)$$

442 It is clear that for any constant k we can pick a number
 443 of features n which render the likelihood of a quadratic behav-
 444 iour infinitesimal. In essence the larger n becomes, the
 445 less likely *Fast-Match* is to exhibit super linear behaviour.

446 The noteworthy part of this result is that for the *retrieval*
 447 variation we can match two images in $O(n)$. Since there are
 448 no large constants involved this makes it possible to rapidly
 449 match very large images. In particular *Fast-Match* is suited
 450 for cases where we are looking to match several large *query*
 451 *images* to one *target image*. In this case we can compute the
 452 features and distances of the *target image* and then match
 453 each *query image* in linear time. In contrast the general
 454 variation still has a complexity of $O(n \log n)$ due to the ne-
 455 cessity of calculating distances online.

456 4. Experimental Setup

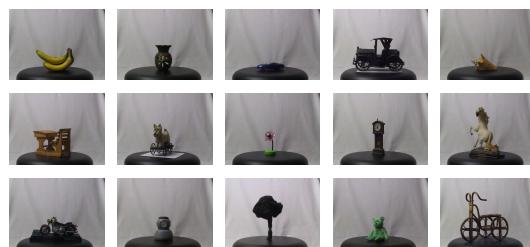
457 We use a dataset of 3024 image pairs featuring various
 458 3D objects at different angles and lighting conditions to
 459 measure matching performance for real life objects when
 460 lighting and perspective changes. To evaluate the per-
 461 formance of *Fast-Match* when only part of (or none of) the
 462 scene is shared between an image pair we conduct exper-
 463 iments on 400 random crops with different degree of image
 464 overlap taken from 4 image pairs, creating a challenging test
 465 set where large part of the images consists of background
 466 clutter similar in nature to the subject being matched. Fi-
 467 nally we time the algorithm on two pairs of images with
 468 high pixel counts to measure performance in terms of speed.

469 We compare the algorithm to the standard *Ratio-Match*
 470 [22] as well as the newer *Mirror-Match* [1]. In addition we
 471 compare the speed of *Fast-Match* to three geometric algo-
 472 rithms; Adaptive Multi-Affing matching (*AMA*) by Suiza et
 473 al. [27], Local shape estimation by Del bimbo et. al (Del
 474 Bimbo) [11] and the Isomatch matching method by Das et
 475 al [10]. Because of the additional time needed to run the
 476 geometric algorithms we do not include them in compari-
 477 ons featuring 3D objects and partially overlapping crops.
 478 However we compare *AMA* to *Fast-Match* and *Ratio-Match*
 479 in a more limited survey featuring the same 4 image pairs
 480 used for the partially overlapping crops since this algorithm
 481 is the fastest according to our results and outperforms *Del*
 482 *Bimbo* according to Suiza et al. [27].

483 *Fast-Match* with three matching algorithms;

484 4.1. Evaluation of Fast-Match on 3D Objects

485 The 3D Objects dataset by Moreels and Perona [24] al-
 486 lows us to experimentally compare matching algorithms
 487 over a large range of object and surface types rotated on a
 488 turnstile and photographed from every angle in increments
 489 of 5 degrees. 15 objects from the dataset are shown in Fig-
 490 ure 3. We use images of 84 different objects photographed
 491 under three different lighting conditions at 12 different an-
 492 gle intervals, conducting experiments with a total of 3024
 493 image pairs. All photos are taken with a consumer camera in
 494 3.1 megapixel resolution.



495 Figure 3: 15 objects from the 3D Objects dataset by Moreels
 496 and Pietro [24].

497 To validate matches, Moreels and Perona propose a
 498 method using epipolar constraints [24, p.266]. According
 499 to their experiments, these constraints are able to identify
 500 true correspondences with an error rate of 2%. We use their
 501 proposed method to generate the ground truth for the eval-
 502 uation of our framework.

503 To compute the total number of possible true corre-
 504 spondences, we take each feature in a *query image* and
 505 count how many of them have a feature in the *target im-*
 506 *age* which would satisfy the epipolar constraints. To avoid
 507 matching non-moving background and foreground objects
 508 we excluded features with no correspondences from being
 509 matched. This was necessary because a relatively small
 510 number of actual true correspondences between folds in the
 511 background material were mistakenly counted as false pos-
 512 itives when evaluated which impacted the precision of test in
 513 particular on cases with few true correspondences. In prac-
 514 tice only few feature points were excluded for this reason.

515 We evaluate all matching algorithms on the 3D Objects
 516 dataset by matching images at different angular intervals.
 517 For each object we pick the *query image* as the image taken
 518 at 10 degrees rotation for calibration stability. We then
 519 match this image with the same object turned an additional
 520 Δ degrees, $\Delta \in \{5, 10, \dots, 60\}$. For every angle interval
 521 we compare images taken under 3 different lighting con-
 522 ditions as provided by the dataset. We include all objects in
 523 the database for which photos at 5 degree angle intervals
 524 are available except for the “Rooster” and “Sponge” objects

540 due to image irregularities. For each angle this means that
 541 a total of 252 image pairs are matched. We calculate the
 542 result these pairs using a weighted average based on the num-
 543 ber of actual correspondences for each pairs. This ensures
 544 that image pairs that are more difficult to match are not con-
 545 tributing less to the end result because less actual matches
 546 would be found.
 547

4.2. Evaluation of Fast-Match on partially overlap- ping scenes



551 Figure 4: Four image pairs from the Graffiti dataset by
 552 Mikolajczyk et al. [23]. From left to right the images are
 553 ‘Boat’, ‘Graf’, ‘Bark’ and ‘Wall’.
 554

555 To evaluate how well *Fast-Match* handles partially over-
 556 lapping scenes we make use of four images from the Graffiti
 557 dataset by Mikolajczyk et al. [23] as displayed in Figure 4
 558 all featuring challenging perspective change. For each of
 559 the four image pairs we generate a hundred pairs of smaller
 560 crops each crop randomly positioned within the original im-
 561 age. This process ensures that image pairs that does not
 562 contain any overlap in scene are still visually similar, making
 563 it more challenging for a matching algorithm to avoid
 564 false positives, i.e. matching local image features that are
 565 not true correspondences.
 566

567 In practice the amount of overlap between pairs in a test
 568 set will depend on the overlap and viewpoint change in the
 569 source image pair. To give a rough idea, Table 1 shows the
 570 overlap of patch pairs created from images 1 and 3 of the
 571 *Graf* image.
 572

573 Table 1: Overlap in the set of 100 patch pairs created from
 574 the ‘Graf’ image pair (Figure 4).
 575

Amount of overlap:	0%	< 50%	> 50%
Number of patch pairs:	21	54	25

576 Given a potential match between two pixels p_1 and p_2 ,
 577 $m = (p_1, p_2)$, and a homography H relating the two images
 578 I_1 and I_2 , we can calculate if m is an inlier by checking if
 579 the two points satisfy the following criteria:
 580

$$|H p_1 - p_2| + |H^{-1} p_2 - p_1| < d_{\max}$$

581 That is, the distance between p_1 translated to I_2 and p_2 plus
 582

594 the distance between p_2 translated to I_1 should be less than
 595 a certain threshold (we use $d_{\max} = 5$ pixels here).
 596

597 As with the 3D Objects we accumulate results across the
 598 hundred image crop pairs by using an average weighted by
 599 the amount of true correspondences between each cropped
 600 pair. This ensures that difficult to match and trivial image
 601 pairs have an equal impact on the evaluation result despite
 602 the difficult image pairs yielding far less actual matches.
 603

4.3. Configuration of Fast-Match

604 The central parameters of *Fast-Match* are the confidence
 605 threshold τ for selecting seed matches and the final confi-
 606 dence threshold applied to the total set of matches. For
 607 the experiments we let $\tau = 0.9$ and created precision/recall
 608 plots by varying the confidence threshold over the final set
 609 of matches.
 610

611 To achieve a good balance between speed and robustness
 612 we let the region in which we extract features be a square
 613 window with a side length of 90 pixels. From empirical
 614 tests we found that this number works well with the SIFT
 615 feature descriptor, but for other descriptors different num-
 616 bers might lead to better performance. The window is split
 617 in to nine smaller squares (see also Figure 2). When a seed
 618 match falls in any of these squares we match only the fea-
 619 tures within the smaller square. We let both the regions and
 620 the smaller squares overlap each other at all edges with 25
 621 pixels in order to capture feature points lying close to an
 622 edge. For the image we have already cached we find all
 623 features within a radius of 50 pixels of the seed match.
 624

5. Results

625 Figure 5 shows the performance of the different match-
 626 ing methods in our proposed framework for 12 increas-
 627 ingly bigger angle differences. The results are shown in
 628 a precision-recall plot to make it easy to compare per-
 629 formance in terms of precision at similar levels of recall.
 630 For each plot we show the accumulated results on all 3D
 631 objects, weighted by the number of possible true corre-
 632 spondences for the individual objects. This ensures that each
 633 object contributes equally to the final result, despite some
 634 objects resulting in disproportionately more matches than oth-
 635 others.
 636

637 At low angular differences all algorithms show similar
 638 performance at low recall, however *Fast-Match* is superior
 639 to *Ratio-Match* and *Mirror-Match* at higher recall, more
 640 than doubling the precision at similar recall levels. This
 641 picture remains the same at higher angular differences, but
 642 with a higher performance gap between *Fast-Match* and the
 643 other algorithms at low recall while more than doubling pre-
 644 cision at higher recalls. For image pairs with angular differ-
 645 ences of more than 40° the performance of all algorithms
 646 declines, with *Fast-Match* still coming out on top.
 647

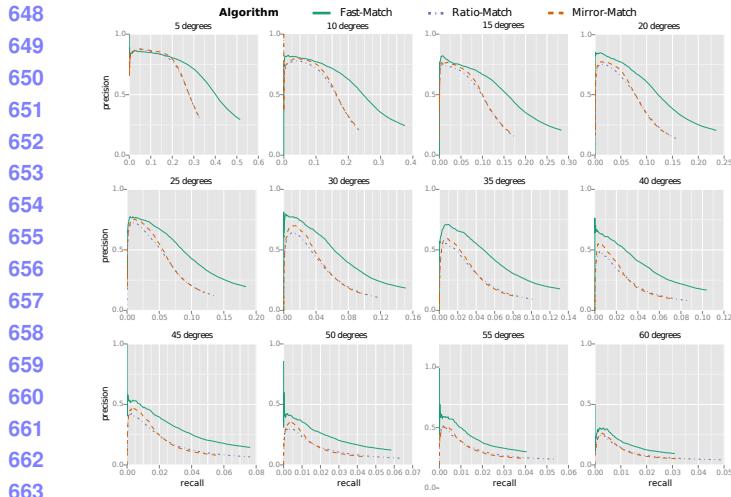


Figure 5: Results for the 3D Objects dataset. Each plot contains the average precision and recall over 84 objects photographed under 3 different lighting conditions weighted over the number of true correspondences for each pair.

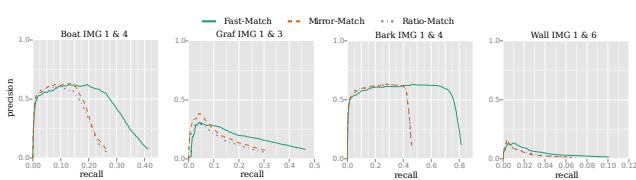


Figure 6: Results for the Graffiti dataset. Each plot contains the average result over 100 semi overlapping image patches based on the same two images weighted over the number of true correspondences for each pair.

The high recall rates of *Fast-Match* are partially explained by the fact that we can allow our confidence thresholds to be more lenient when we already know that we are likely to find true correspondences in the regions that we are matching. However, lenient thresholds can easily impact the matching speed, so in practical situations we usually have to choose between matching fast and precisely, or more slowly with higher recall.

In Figure 6 we see the performance of the three different matching methods over 400 crop pairs based on four image pairs. As for the 3D Objects the results are shown in a precision-recall plot weighted by the number of possible true correspondences for the individual crop pairs.

At low recall we see a similar performance across all algorithms, but as the algorithms return more matches, *Fast-Match* shows drastically better precision at similar recall levels than *Mirror-Match* and *Ratio-Match*. While this difference is more noticeable in the case of the ‘Bark’ and

¹As found here: http://ranger.uta.edu/~gianluca/feature_matching/

Table 2: *Fast-Match* and *Ratio-Match* compared with *AMA* on 4 image pairs from the Graffiti dataset by Mikolajczyk et al. [23] as pictured in Figure 4. *Fast-Match* and *Ratio-Match* are evaluated with a threshold of $\tau = 0.6$ while *AMA* is used with default parameters according to the reference implementation¹. Precision is listed first with number of correct matches in parenthesis.

		<i>Fast-Match</i>	<i>Ratio-Match</i>	<i>AMA</i>
Bark 1-4:	0.99	(1091)	0.99	(645)
Boat 1-4:	0.99	(672)	0.97	(455)
Graf 1-3:	0.67	(192)	0.72	(141)
Wall 1-6:	0.00	(0)	0.00	(0)

‘Boat’ images, the ‘Graf’ and ‘Wall’ also see a drastic improvement in recall but as precision is lower in general for these two images the effect is less visible.

For all four image pairs we see a precision dip at very low recall which is explained by the presence of image crop pairs with no or little overlap. When matching these pairs we might often find matches even though there are no or few possible true correspondences which in effect negatively impacts the results at very low recall rates.

Table 2 lists the precision and number of correct matches from matching 4 image pairs as seen in Figure 4. The geometric matching method *AMA* yields better results than *Fast-Match* and *Ratio-Match* for the more challenging ‘Graf’ image pair while all algorithms yield no correct matches on the ‘Wall’ pair. For the more trivial ‘Boat’ and ‘Bark’ images we find *Fast-Match* performing similar to *Ratio-Match* in terms of precision while returning markedly more correct matches.

Figure 7 demonstrate the speed of *Fast-Match* on image sizes as is typically produced by modern consumer cameras and camera phones. We compare *Fast-Match* with different variations of *Ratio-Match* and three geometric matching algorithms on two image pairs, one with two images of 10 Mpixels and another where the same image pair has been scaled to 3 Mpixels. The differences between *Fast-Match* and *Ratio-Match* are largest for larger images. The retrieval variant of *Fast-Match* matches the image pairs in around one second while a retrieval variant of *Ratio-Match* with precomputed features spends eight seconds on the same task. This result is roughly equal to the time the general variation of *Fast-Match* spends matching the two images without any pre-computations. For the nearest neighbor search we note that for large images there are substantial speed gains to be had by choosing the right algorithm for finding nearest neighbors, with Muja et al.’s Flann matcher being vastly superior to brute force and KGraph. For the 3 megapixel picture the pattern repeats itself, although with smaller margins separating *Fast-Match* and *Ratio-Match*.

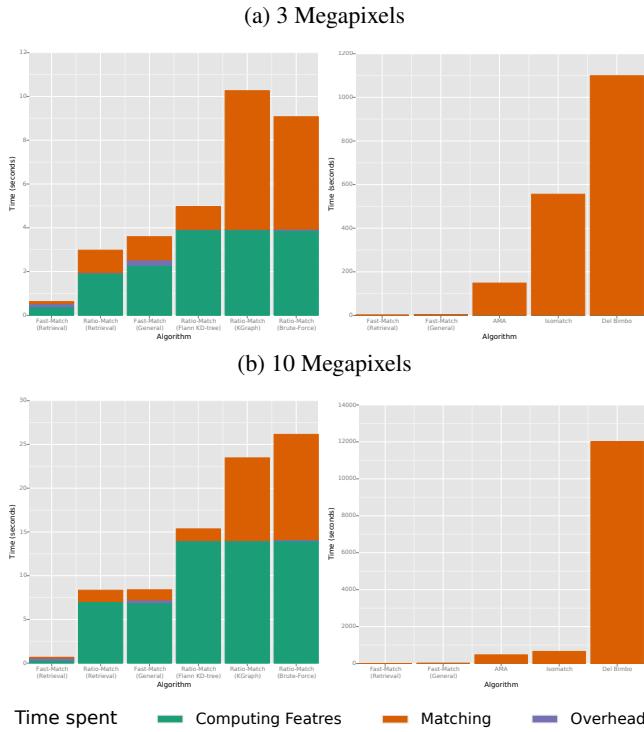


Figure 7: The duration in seconds the two versions of *Fast-Match* compared with *Ratio-Match* and the geometric matching algorithms *AMA*, *Isomatch* and *Del Bimbo*. The two top plots shows the duration in seconds for matching two 3 MPixel images. The bottom plots shows the duration in seconds for matching two 10 Mpixel images. For the plots on the left, the first two results of each plot presume a precached image while the rest do not require knowing any inputs beforehand. The two plots on the right shows the timings for the precached and general version of *Fast-Match* identical to the plots on the left but this time compared with three geometric matching algorithms. Note that the y-scale is different on all plots. The image pairs used can be seen in Figure 1

When images are smaller, computations like obtaining the initial seed matches are comparatively more costly.

In the plots featuring geometric matching algorithms we compare *Fast-Match* with three matching algorithms; Adaptive Multi-Affing matching (*AMA*) by Suiza et al. [27], Local shape estimation by Del bimbo et. al (*Del Bimbo*) [11] and the Isomatch matching method by Das et al [10]. For both the 3 Mpixel and 10 Mpixel images the differences between *Fast-Match* and the geometric algorithms in terms of speed is vast. For the 3 Mpixel images the general *Fast-Match* method is roughly 40 times faster than the closest competitor, *AMA* while it is more then 50 times as fast as *AMA* for the 10 Mpixel images. While both *Isomatch* and *AMA* are slow they scale well with the increased image

sizes. *Del Bimbo* on the other hand spends more than an hour finding matches for the 10 Mpixel image pairs.

Figure 1 shows the result from matching the pair of three megapixel images using *Fast-Match* and *Ratio-Match*. *Fast-Match* ends up only matching the part of the images that are likely to have correspondences. This highlights both a strength and a weakness of *Fast-Match*. It is this myopic matching approach that is responsible for *Fast-Match*' speed and precision; on the other hand, *Fast-Match* is not the best candidate for images that are almost identical, because its step by step approach incurs too much overhead to compete in speed with *Ratio-Match* in those cases.

6. Conclusion

We have introduced *Fast-Match* in two forms, a general and a retrieval variation. The retrieval variation assumes that we know one of the images that we are matching ahead of time in order to pre-cache feature points. The general variation has no assumptions and performs the pre-caching of feature points online instead. We have proven that while the general variation has a computational complexity of $O(n \log n)$ with n being the number of feature points, the retrieval variation will on average run in $O(n)$. From experiments on large images we have shown that *Fast-Match* can be a magnitude faster than *Ratio-Match* while providing comparable results. We further compared *Fast-Match* to *Ratio-Match* and *Mirror-Match* using images of rotated 3D objects to demonstrate that *Fast-Match* outperforms the other algorithms significantly over 3024 image pairs and in most cases doubles the matching precision at similar recall rates. We show that these results hold true for matching cases featuring partial or no overlap by comparing *Fast-Match* on 400 crop pairs of planar scenes. Finally we compared *Fast-Match* with the geometric matching method *AMA* and note that while *AMA* can yield superior matches for challenging images. However we measured *Fast-Match* to be more to 40 times faster than *AMA* on larger images while performing much better on image pairs that are easier to match. *Fast-Match* is particularly applicable in cases where we are interested in matching one image to multiple large images, in which case we can quickly pre-cache the image and use this data for each of the other images.

References

- [1] J. T. Arnfred, S. Winkler, and S. Süsstrunk. Mirror Match: Reliable feature point matching without geometric constraints. In *Proc. IAPR Asian Conference on Pattern Recognition (ACPR)*, pages 256–260, 2013.
- [2] A. Baumberg. Reliable feature matching across widely separated views. In *Proc. Conference on Computer Vision and Pattern Recognition (CVPR)*, volume 1, pages 774–781, 2000.

810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863

- 864 [3] H. Bay, T. Tuytelaars, and L. Van Gool. SURF:
865 Speeded up robust features. In *Proc. European Conference on Computer Vision (ECCV)*, pages 404–417,
866 2006.
867 [4] J. S. Beis and D. G. Lowe. Shape indexing using approximate nearest-neighbour search in high-dimensional spaces. In *Proc. Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 1000–1006. IEEE, 1997.
868 [5] M. Brown, R. Szeliski, and S. Winder. Multi-image matching using multi-scale oriented patches. In *Proc. Conference on Computer Vision and Pattern Recognition (CVPR)*, volume 1, pages 510–517, 2005.
869 [6] M. Calonder, V. Lepetit, C. Strecha, and P. Fua. BRIEF: binary robust independent elementary features. In *Proc. European Conference on Computer Vision (ECCV)*, pages 778–792, 2010.
870 [7] H.-Y. Chen, Y.-Y. Lin, and B.-Y. Chen. Robust feature matching with alternate hough and inverted hough transforms. In *Proc. Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 2762–2769. IEEE, 2013.
871 [8] M. Cho, J. Lee, and K. M. Lee. Feature correspondence and deformable object matching via agglomerative correspondence clustering. In *Proc. International Conference on Computer Vision (ICCV)*, pages 1280–1287. IEEE, 2009.
872 [9] M. Cho, J. Lee, and K. M. Lee. Reweighted random walks for graph matching. In *Proc. European Conference on Computer Vision (ECCV)*, pages 492–505. Springer, 2010.
873 [10] M. Das, J. Farmer, A. Gallagher, and A. Loui. Event-based location matching for consumer image collections. In *Proc. Content-based Image and Video Retrieval (CIVR)*, pages 339–348, Niagara Falls, ON, 2008.
874 [11] A. Del Bimbo, F. Franco, and F. Pernici. Local shape estimation from a single keypoint. In *Proc. Computer Vision and Pattern Recognition Workshops(CVPRW)*, pages 23–28. IEEE, 2010.
875 [12] R. Deriche, Z. Zhang, Q.-T. Luong, and O. Faugeras. Robust recovery of the epipolar geometry for an uncalibrated stereo rig. In *Proc. European Conference on Computer Vision (ECCV)*, pages 567–576, 1994.
876 [13] W. Dong, C. Moses, and K. Li. Efficient k-nearest neighbor graph construction for generic similarity measures. In *Proc. International Conference on World Wide Web (IW3C2)*, pages 577–586. ACM, 2011.
877 [14] M. A. Fischler and R. C. Bolles. Random sample consensus: a paradigm for model fitting with applications
878 to image analysis and automated cartography. *Communications of the ACM*, 24(6):381–395, 1981.
879 [15] J. Heinly, E. Dunn, and J.-M. Frahm. Comparative evaluation of binary features. In *Proc. European Conference on Computer Vision (ECCV)*, pages 759–773, 2012.
880 [16] Y. Ke and R. Sukthankar. PCA-SIFT: A more distinctive representation for local image descriptors. In *Proc. Conference on Computer Vision and Pattern Recognition (CVPR)*, volume 2, pages II–506. IEEE, 2004.
881 [17] J. Kim, O. Choi, and I. S. Kweon. Efficient feature tracking for scene recognition using angular and scale constraints. In *Proc. IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 4086–4091, Nice, France, 2008.
882 [18] M. Leordeanu and M. Hebert. A spectral technique for correspondence problems using pairwise constraints. In *Proc. Conference on Computer Vision and Pattern Recognition (CVPR)*, volume 2, pages 1482–1489, San Diego, CA, 2005.
883 [19] S. Leutenegger, M. Chli, and R. Y. Siegwart. BRISK: Binary robust invariant scalable keypoints. In *Proc. International Conference on Computer Vision (ICCV)*, pages 2548–2555. IEEE, 2011.
884 [20] J. Li and N. M. Allinson. A comprehensive review of current local features for computer vision. *Neurocomputing*, 71(10):1771–1787, 2008.
885 [21] D. G. Lowe. Object recognition from local scale-invariant features. In *Proc. International Conference on Computer Vision (ICCV)*, volume 2, pages 1150–1157. Ieee, 1999.
886 [22] D. G. Lowe. Distinctive image features from scale-invariant keypoints. *International Journal on Computer Vision*, 60(2):91–110, 2004.
887 [23] K. Mikolajczyk and C. Schmid. A performance evaluation of local descriptors. *IEEE Trans. Pattern Analysis and Machine Intelligence*, 27(10):1615–1630, 2005.
888 [24] P. Moreels and P. Perona. Evaluation of features detectors and descriptors based on 3D objects. *International Journal on Computer Vision*, 73(3):263–284, 2007.
889 [25] M. Muja and D. G. Lowe. Fast approximate nearest neighbors with automatic algorithm configuration. In *Proc. International Joint Conference on Computer Vision, Imaging and Computer Graphics Theory and Applications (VISAPP)*, pages 331–340, 2009.
890 [26] Y. Pang, M. Shang, Y. Yuan, and J. Pan. Scale invariant image matching using triplewise constraint and weighted voting. *Neurocomputing*, 83:64–71, 2012.
891 [27] M. Pollefeys, M. van Gool, and R. Koch. Multi-camera
892 [28] [29] [30] [31] [32] [33] [34] [35] [36] [37] [38] [39] [40] [41] [42] [43] [44] [45] [46] [47] [48] [49] [50] [51] [52] [53] [54] [55] [56] [57] [58] [59] [60] [61] [62] [63] [64] [65] [66] [67] [68] [69] [70] [71] [72] [73] [74] [75] [76] [77] [78] [79] [80] [81] [82] [83] [84] [85] [86] [87] [88] [89] [90] [91] [92] [93] [94] [95] [96] [97] [98] [99] [100] [101] [102] [103] [104] [105] [106] [107] [108] [109] [110] [111] [112] [113] [114] [115] [116] [117] [118] [119] [120] [121] [122] [123] [124] [125] [126] [127] [128] [129] [130] [131] [132] [133] [134] [135] [136] [137] [138] [139] [140] [141] [142] [143] [144] [145] [146] [147] [148] [149] [150] [151] [152] [153] [154] [155] [156] [157] [158] [159] [160] [161] [162] [163] [164] [165] [166] [167] [168] [169] [170] [171] [172] [173] [174] [175] [176] [177] [178] [179] [180] [181] [182] [183] [184] [185] [186] [187] [188] [189] [190] [191] [192] [193] [194] [195] [196] [197] [198] [199] [200] [201] [202] [203] [204] [205] [206] [207] [208] [209] [210] [211] [212] [213] [214] [215] [216] [217] [218] [219] [220] [221] [222] [223] [224] [225] [226] [227] [228] [229] [230] [231] [232] [233] [234] [235] [236] [237] [238] [239] [240] [241] [242] [243] [244] [245] [246] [247] [248] [249] [250] [251] [252] [253] [254] [255] [256] [257] [258] [259] [260] [261] [262] [263] [264] [265] [266] [267] [268] [269] [270] [271] [272] [273] [274] [275] [276] [277] [278] [279] [280] [281] [282] [283] [284] [285] [286] [287] [288] [289] [290] [291] [292] [293] [294] [295] [296] [297] [298] [299] [300] [301] [302] [303] [304] [305] [306] [307] [308] [309] [310] [311] [312] [313] [314] [315] [316] [317] [318] [319] [320] [321] [322] [323] [324] [325] [326] [327] [328] [329] [330] [331] [332] [333] [334] [335] [336] [337] [338] [339] [340] [341] [342] [343] [344] [345] [346] [347] [348] [349] [350] [351] [352] [353] [354] [355] [356] [357] [358] [359] [360] [361] [362] [363] [364] [365] [366] [367] [368] [369] [370] [371] [372] [373] [374] [375] [376] [377] [378] [379] [380] [381] [382] [383] [384] [385] [386] [387] [388] [389] [390] [391] [392] [393] [394] [395] [396] [397] [398] [399] [400] [401] [402] [403] [404] [405] [406] [407] [408] [409] [410] [411] [412] [413] [414] [415] [416] [417] [418] [419] [420] [421] [422] [423] [424] [425] [426] [427] [428] [429] [430] [431] [432] [433] [434] [435] [436] [437] [438] [439] [440] [441] [442] [443] [444] [445] [446] [447] [448] [449] [450] [451] [452] [453] [454] [455] [456] [457] [458] [459] [460] [461] [462] [463] [464] [465] [466] [467] [468] [469] [470] [471] [472] [473] [474] [475] [476] [477] [478] [479] [480] [481] [482] [483] [484] [485] [486] [487] [488] [489] [490] [491] [492] [493] [494] [495] [496] [497] [498] [499] [500] [501] [502] [503] [504] [505] [506] [507] [508] [509] [510] [511] [512] [513] [514] [515] [516] [517] [518] [519] [520] [521] [522] [523] [524] [525] [526] [527] [528] [529] [530] [531] [532] [533] [534] [535] [536] [537] [538] [539] [540] [541] [542] [543] [544] [545] [546] [547] [548] [549] [550] [551] [552] [553] [554] [555] [556] [557] [558] [559] [550] [551] [552] [553] [554] [555] [556] [557] [558] [559] [560] [561] [562] [563] [564] [565] [566] [567] [568] [569] [570] [571] [572] [573] [574] [575] [576] [577] [578] [579] [580] [581] [582] [583] [584] [585] [586] [587] [588] [589] [580] [581] [582] [583] [584] [585] [586] [587] [588] [589] [590] [591] [592] [593] [594] [595] [596] [597] [598] [599] [590] [591] [592] [593] [594] [595] [596] [597] [598] [599] [600] [601] [602] [603] [604] [605] [606] [607] [608] [609] [610] [611] [612] [613] [614] [615] [616] [617] [618] [619] [610] [611] [612] [613] [614] [615] [616] [617] [618] [619] [620] [621] [622] [623] [624] [625] [626] [627] [628] [629] [630] [631] [632] [633] [634] [635] [636] [637] [638] [639] [640] [641] [642] [643] [644] [645] [646] [647] [648] [649] [650] [651] [652] [653] [654] [655] [656] [657] [658] [659] [660] [661] [662] [663] [664] [665] [666] [667] [668] [669] [660] [661] [662] [663] [664] [665] [666] [667] [668] [669] [670] [671] [672] [673] [674] [675] [676] [677] [678] [679] [680] [681] [682] [683] [684] [685] [686] [687] [688] [689] [690] [691] [692] [693] [694] [695] [696] [697] [698] [699] [690] [691] [692] [693] [694] [695] [696] [697] [698] [699] [700] [701] [702] [703] [704] [705] [706] [707] [708] [709] [710] [711] [712] [713] [714] [715] [716] [717] [718] [719] [720] [721] [722] [723] [724] [725] [726] [727] [728] [729] [723] [724] [725] [726] [727] [728] [729] [730] [731] [732] [733] [734] [735] [736] [737] [738] [739] [730] [731] [732] [733] [734] [735] [736] [737] [738] [739] [740] [741] [742] [743] [744] [745] [746] [747] [748] [749] [745] [746] [747] [748] [749] [750] [751] [752] [753] [754] [755] [756] [757] [758] [759] [750] [751] [752] [753] [754] [755] [756] [757] [758] [759] [760] [761] [762] [763] [764] [765] [766] [767] [768] [769] [760] [761] [762] [763] [764] [765] [766] [767] [768] [769] [770] [771] [772] [773] [774] [775] [776] [777] [778] [779] [770] [771] [772] [773] [774] [775] [776] [777] [778] [779] [780] [781] [782] [783] [784] [785] [786] [787] [788] [789] [780] [781] [782] [783] [784] [785] [786] [787] [788] [789] [790] [791] [792] [793] [794] [795] [796] [797] [798] [799] [790] [791] [792] [793] [794] [795] [796] [797] [798] [799] [800] [801] [802] [803] [804] [805] [806] [807] [808] [809] [810] [811] [812] [813] [814] [815] [816] [817] [818] [819] [810] [811] [812] [813] [814] [815] [816] [817] [818] [819] [820] [821] [822] [823] [824] [825] [826] [827] [828] [829] [820] [821] [822] [823] [824] [825] [826] [827] [828] [829] [830] [831] [832] [833] [834] [835] [836] [837] [838] [839] [830] [831] [832] [833] [834] [835] [836] [837] [838] [839] [840] [841] [842] [843] [844] [845] [846] [847] [848] [849] [840] [841] [842] [843] [844] [845] [846] [847] [848] [849] [850] [851] [852] [853] [854] [855] [856] [857] [858] [859] [850] [851] [852] [853] [854] [855] [856] [857] [858] [859] [860] [861] [862] [863] [864] [865] [866] [867] [868] [869] [860] [861] [862] [863] [864] [865] [866] [867] [868] [869] [870] [871] [872] [873] [874] [875] [876] [877] [878] [879] [870] [871] [872] [873] [874] [875] [876] [877] [878] [879] [880] [881] [882] [883] [884] [885] [886] [887] [888] [889] [880] [881] [882] [883] [884] [885] [886] [887] [888] [889] [890] [891] [892] [893] [894] [895] [896] [897] [898] [899] [890] [891] [892] [893] [894] [895] [896] [897] [898] [899] [900] [901] [902] [903] [904] [905] [906] [907] [908] [909] [900] [901] [902] [903] [904] [905] [906] [907] [908] [909] [910] [911] [912] [913] [914] [915] [916] [917] [918] [919] [910] [911] [912] [913] [914] [915] [916] [917] [918] [919] [920] [921] [922] [923] [924] [925] [926] [927] [928] [929] [930] [931] [932] [933] [934] [935] [936] [937] [938] [939] [940] [941] [942] [943] [944] [945] [946] [947] [948] [949] [950] [951] [952] [953] [954] [955] [956] [957] [958] [959] [960] [961] [962] [963] [964] [965] [966] [967] [968] [969] [970] [971] [972]

- 972 [27] G. A. Puerto Souza, M. Adibi, J. A. Cadeddu, and 1026
973 G. L. Mariottini. Adaptive multi-affine (ama) feature- 1027
974 matching algorithm and its application to minimally- 1028
975 invasive surgery images. In *Proc. Intelligent Robots 1029*
976 and Systems (IROS), pages 2371–2376. IEEE, 2011. 1030
977
978 [28] J. Rabin, J. Delon, and Y. Gousseau. A statistical 1031
979 approach to the matching of local features. *SIAM Journal 1032*
980 on Imaging Sciences, 2(3):931–958, 2009. 1033
981 [29] C. Schmid and R. Mohr. Local grayvalue invariants 1034
982 for image retrieval. *IEEE Transactions on Pattern 1035*
983 Analysis and Machine Intelligence, 19(5):530– 1036
984 535, 1997. 1037
985 [30] R. Szeliski. *Computer Vision: Algorithms and 1038*
986 Applications. Springer, 2010. 1039
987 [31] P. H. S. Torr and A. Zisserman. MLESAC: A new 1040
988 robust estimator with application to estimating image 1041
989 geometry. *Computer Vision and Image Understanding*, 1042
990 78(1):138–156, 2000. 1043
991 [32] L. Torresani, V. Kolmogorov, and C. Rother. Feature 1044
992 correspondence via graph matching: Models and 1045
993 global optimization. In *Proc. European Conference on 1046*
994 Computer Vision (ECCV), pages 596–609. Springer, 1047
995 2008. 1048
996 [33] L. Wu, Y. Niu, H. Zhang, H. Zhu, and L. Shen. Robust 1049
997 feature point matching based on local feature groups 1050
998 (LFGs) and relative spatial configuration. *Journal of 1051*
1000 Computational Information Systems, 7(9):3235–3244, 1052
1001 2011. 1053
1002 [34] J. Yarkony, C. Fowlkes, and A. Ihler. Covering trees 1054
1003 and lower-bounds on quadratic assignment. In *Proc. 1055*
1004 Conference on Computer Vision and Pattern Recog- 1056
1005 nition (CVPR), pages 887–894. IEEE, 2010. 1057
1006 [35] Y. Yuan, Y. Pang, K. Wang, and M. Shang. Efficient 1058
1007 image matching using weighted voting. *Pattern 1059*
1008 Recognition Letters, 33(4):471–475, 2012. 1060
1009
1010
1011
1012
1013
1014
1015
1016
1017
1018
1019
1020
1021
1022
1023
1024
1025