

# Fast-Match: Fast and Robust Feature Matching on Large Images

Anonymous ACCV 2014 submission

Paper ID 325

**Abstract.** Consumer digital cameras and camera phones produce images that often exceed 10 megapixels. Yet computing and matching local features in images of this size can easily take more than twenty seconds using fast matching algorithms. This is much too slow for interactive applications and much too expensive for large scale image operations. We introduce *Fast-Match*, an algorithm designed to swiftly match large images without compromising on matching precision or recall. *Fast-Match* derives its speed from only computing features in those parts of the image that can be confidently matched. We show that *Fast-Match* is an order of magnitude faster than *Ratio-Match*, yet often doubles the precision on difficult to match image pairs at equal recall. In addition we prove that when one image is known in advance, *Fast-Match* can achieve linear performance  $O(n)$  in the number of feature points  $n$ .

## 1 Introduction

Whenever we match local image features we are faced with a choice between performance and precision. On one hand SIFT features proposed by Lowe [1] have shown again and again to compare favorably to other local image descriptors, especially under unconstrained conditions such as perspective change with non-planar scenes [2–4]. On the other, SIFT keypoints and descriptors are slow to compute, the main *raison d'être* for the introduction of various alternative local image features. In many applications of computer vision we would like the increase the computational performance in order to work on larger images, bigger image sets, at faster frame rates or with more limited hardware, but we cannot give up the additional precision that SIFT affords us over other local features without negative effects for our application.

In this paper we introduce a matching algorithm designed to match features between two images only in image areas that are likely to correspond. This approach is much faster than traditional methods because there is no need to compute descriptors for areas in the image that are not matched. We provide two variations of the algorithm: The *general* variation functions as a traditional matching algorithm and matches two unknown images albeit faster and more robustly than conventional methods. The *retrieval* variation on the other hand assumes that we know one of the images we intend to match beforehand; under

CONFIDENTIAL REVIEW COPY. DO NOT DISTRIBUTE.

2 ACCV-14 submission ID 325

045 this assumption, it can be a magnitude faster than existing matching methods  
 046 without compromising on accuracy. We will unimaginatively refer to the pro-  
 047 posed algorithm as *Fast-Match* in this paper and make clear from the context  
 048 whether this refers to the *retrieval* or *general* variation.

049 The problem that *Fast-Match* attempts to solve is two-fold. By matching only  
 050 image areas that are likely to correspond we hope to improve accuracy by entirely  
 051 ignoring parts of the images that would otherwise likely be a source of incorrect  
 052 correspondences. At the same time, this enables us to improve computational  
 053 speed by not having to compute keypoints and descriptors for large parts of the  
 054 images and at the same time reducing the amount of feature points we need to  
 055 match.

056 *Fast-Match* makes use of an angular assumption to efficiently find new matches  
 057 in the geometric neighbourhood of already confirmed matches. However this con-  
 058 straint is only applied *locally* to increase the number of matches, making *Fast-*  
 059 *Match* robust to outliers. In addition *Fast-Match* does not rely on an initial set  
 060 of matches and derives much of its speed from the fact that not even an initial  
 061 set of local image features is required.

062 This paper is structured as follows: Section 2 discuss related work. Section 3  
 063 introduces the *Fast-Match* algorithm. Section 4 outlines the experimental setup.  
 064 In Section 5 we discuss the results before concluding in Section 6.

## 067 2 Related work

069 As noted, the cost of computing SIFT keypoints and descriptors is addressed by  
 070 other local image features such as SURF [5], BRIEF [6], or BRISK [7], just to  
 071 mention a few. Similarly efforts have been made to improve SIFT itself such as  
 072 PCA-SIFT [8] and GLOH [2]. Both apply PCA to reduce descriptor lengths and  
 073 improve distinctiveness, but neither have been shown to consistently outperform  
 074 SIFT [2, 9].

075 Efforts to reduce the computational costs of finding nearest neighbours to  
 076 feature points have largely focused on metric trees. Naively the set of nearest  
 077 neighbours between features in two images can computed by brute force in  $O(n^2)$ ,  
 078 where  $n$  is the total number of feature points in the two images (a typical three  
 079 megapixel image may contain anywhere from 500 to 5000 feature points). When  
 080 SIFT was originally published, Lowe proposed using the Best-Bin-First method  
 081 to approximately search for nearest neighbours [10, 11]. This reduces the com-  
 082 putational complexity to  $O(n \log n)$ , but for smaller images even approximate  
 083 metric trees are hard pressed to compete with brute force due to the high di-  
 084 mensionality of SIFT and the costs incurred with constructing the metric tree.  
 085 Later work by Muja and Lowe [12] focused on improving approximate nearest  
 086 neighbour searches by using several KD-Trees simultaneously while optimizing  
 087 the tree structure using K-Means to cluster similar features. Recent work on knn-  
 088 graphs shows a lot of promise for high dimensional cases [13]. We later review  
 089 these improvements and their effect on efficiently matching large scale images.

090 A wealth of matching methods have focused instead on increasing the efficiency of local feature matching. *Fast-Match* builds upon the foundation of  
091 *Ratio-Match* introduced originally by Deriche et al. [14] and Baumberg [15] even  
092 though Lowe is usually credited for introducing it [1]. They both propose using  
093 the ratio of the similarity of the best to second best match of a given point to evaluate  
094 how unique it is. Their finding has later been tested by several independent  
095 teams, all concluding that thresholding based on this ratio is generally superior  
096 to thresholding based on similarity or returning all nearest neighbours [1–3, 16].  
097

098 Brown and Lowe [17] extend ratio match to deal with a set of images by using  
099 not the ratio of the best and second best match, but the average ratio of the best  
100 and the average of second best matches across a set of images. Rabin et al. [16]  
101 try to enhance descriptor matching by looking at the statistical distribution  
102 of local features in the matched images, and only return a match when such a  
103 correspondence would not occur by mere chance. Finally, Arnfred et al. generalize  
104 *Ratio-Match* to make use of both of the matched images to provide a more  
105 accurate evaluation of the uniqueness of a match [18].

106 Another inspiration for the design of *Fast-Match* is *Patch-Match* as introduced  
107 by Barnes et al. [19]. Like *Fast-Match*, *Patch-Match* is an iterative algorithm  
108 for fast image matching, but unlike *Fast-Match*, which uses sparse local  
109 image features, *Patch-Match* is designed for dense image features. It works by  
110 randomly creating a set of matches from both images and then iteratively improving  
111 it.

112 For sparse local image features many solutions have combined *Ratio-Match*  
113 with various geometric constraints to improve matching. These constraints are  
114 based on assumptions regarding the transformation between the *query* and *target images*.  
115 A commonly used example is *RANSAC* applied to feature matching where matches are  
116 chosen from a pool of candidates according to how well they approximate a global  
117 epipolar geometry [20–22]. Similar global angular and distance constraints can be  
118 used to filter a set of matches as shown in [23, 24]. Finally the problem of feature  
119 matching can be modelled as a graph matching problem where each feature is a vertex,  
120 and edge values correspond to a geometric relation between two features as demonstrated  
121 in [25–27]. For cases that require more flexibility such as scenes with moving elements and non-planar  
122 perspective change pairwise constraints provide an alternative to the more rigid global  
123 constraints. Introduced by Leordeanu and Herbert [28], pairwise constraints work by  
124 applying a geometric constraint across matches on a pairwise basis, attempting  
125 to minimize the pairwise error. This approach has later been adopted by Pang et  
126 al. but made scale invariant and more efficient by using a voting scheme [29, 30].  
127 Similarly we can cluster matches together based on a geometric constraint and  
128 treat each cluster as a separate case [31, 32]. An interesting application of this  
129 principle is demonstrated by Chen et al. who iteratively use a Hough transform  
130 to cluster matches using angular constraints [33].

132 While geometric constraints have been shown to work well, they are often  
133 susceptible to outliers and tend to be computationally demanding. All of the  
134 above geometric methods require a set of initial matches usually provided by



Fig. 1: The 50 best matches by the proposed *Fast-Match* algorithm (left) and the *Ratio-Match* algorithm [1] (right). The lines between the image pairs denote a match as proposed by the algorithm. For *Fast-Match* the areas where local features were computed are marked with blue squares in the right image.

*Ratio-Match* which acts as a lower bound on their running time. In practice even fast graph matching methods such as the covering trees method proposed by Yarkony et al. are two or three magnitudes slower than *Ratio-Match* [26].

### 3 Matching Fast and Slow

In this section we will introduce the fundamentals of *Fast-Match* and motivate the design choices as we go. The algorithm consists of 3 components; finding seeds, finding matches, and exploring for other places where matches might be.

#### 3.1 Considerations

If we set out to design a truly fast matching algorithm, we cannot rely just on optimizing the matching step once the descriptors from both images have been found and computed. Finding and computing descriptors can easily account for 80% of the time spent matching in the case of bigger images (cf. Figure 5). For this reason *Fast-Match* is designed to only compute features for the part of the image we hope to match.

The *Fast-Match* algorithm is demonstrated in Algorithm 1. Given a *query image* and a *target image* that we intend to match and a confidence threshold  $\tau$ , we obtain a set of seed matches from the two images. For each seed match we look at the matched position in the *query* and *target images* and find a set of matches. We save these matches and their confidence scores; for those that pass the confidence threshold  $\tau$ , we obtain another set of seed matches. In this way we iterate until we have no more seed matches and return the matches and their confidence scores. In an intuitive sense  $\tau$  serves as a parameter directing the thoroughness of the algorithm, i.e. how much time we spend matching, while the final precision and recall can be adjusted by thresholding the matches on their confidence scores at the end.

We introduce a *general* and a *retrieval* variation of the algorithm. The latter assumes that we know one of the images that we intend to match ahead of time

180 and can do some off-line computations. For the *general* variation we make those  
 181 computations on the fly instead and do not assume anything to be pre-computed.  
 182

---

**183 Algorithm 1** Fast-Match
 

---

```

185 Require:  $I_{query}, I_{target}$  : images, maxiter  $\in \mathbb{N}$ ,  $\tau \in [0, 1]$ 
186  $M_{seed} \leftarrow \text{seed\_matches}(I_{query}, I_{target})$ 
187  $M_{final} \leftarrow \emptyset$ 
188  $C_{final} \leftarrow \emptyset$ 
189  $M_{seen} \leftarrow \emptyset$ 
190 while  $M_{seed} \neq \emptyset \wedge i < \text{maxiter}$  do
191    $M_{round} \leftarrow \text{get\_matches}(M_{seed})$ 
192    $C_{round} \leftarrow \text{get\_confidence}(M_{round})$ 
193    $M_{seed} \leftarrow \text{get\_seeds}(M_{round} \setminus M_{seen}, C_{round}, \tau)$ 
194    $M_{seen} \leftarrow M_{seen} \cup M_{seed}$ 
195    $M_{final} \leftarrow M_{final} \cup M_{round}$ 
196    $C_{final} \leftarrow C_{final} \cup C_{round}$ 
197 end while
198 return  $M_{final}, C_{final}$ 

```

---

**201 3.2 Initiating Seeds**

203 Depending on the scenario, several strategies can be used to obtain a set of  
 204 initial seed matches. For the case of matching a pair of two images we can get a  
 205 rough group of seed matches by matching the thumbnails of the two images with  
 206 for example *Ratio-Match*. However if we were matching images in a series, such  
 207 as frames from a movie we could instead make use of a subset of the matches  
 208 from the last frame to seed *Fast-Match* on the next. In practice we have found  
 209 it efficient to resize both images to thumbnails of  $300 \times 300$  pixels and use ratio  
 210 match to obtain a set of matches and ratios that we then threshold with the  
 211 confidence threshold  $\tau$  to obtain the initial seed matches.  
 212

**213 3.3 Collecting Matches and Computing Confidence**

214 If a seed match yields a connection between two points  $p_q$  in the *query image* and  
 215  $p_t$  in the *target image*, then we are interested in collecting all matches between  
 216 the regions  $R_q$  and  $R_t$  centered around  $p_q$  and  $p_t$  respectively. From each region  
 217 we can extract a set of feature points between which we try to find a set of  
 218 matches  $M_{qt}$  and a set of confidence scores  $C_{qt}$ .

219 Lowe and others have shown that the distance between two SIFT descriptors  
 220 is much less indicative of a true correspondence than the ratio between the best  
 221 and second best match [1–3,16]. This ratio is more formally defined as follows: If  
 222 we let  $f_q$  be a feature in the *query image* and  $f_t, f_b$  be the two nearest neighbours  
 223

224

225 of  $f_q$  in the *target image* then the ratio  $r$  is defined as follows:

$$226 \\ 227 \quad r = \frac{d(f_q, f_t)}{d(f_q, f_b)}. \\ 228$$

229 Here  $d(f_i, f_j)$  is the distance between the features  $f_i$  and  $f_j$ . For SIFT this is  
 230 the Euclidean distance. The distance between two feature points is greater when  
 231 the feature points resemble each other less. For this reason we have higher con-  
 232 fidence in a match with a low  $r$ -value. Using  $r$  as a measure of match confidence  
 233 presumes that we expect each feature in the *target image* to have at most one  
 234 true correspondence in the *target image*. Intuitively if we try to find a match  
 235 for a feature  $f_i$  in an image that does not have any true correspondences, then  
 236 we would expect the two closest neighbours to be roughly equally well matched  
 237 with  $f_i$ . For this reason we attribute high confidence to matches where the clos-  
 238 est neighbour is dramatically closer to  $f_i$  than the second closest neighbour and  
 239 discard the rest.

240 We are faced with a problem when applying this technique to obtain the set  
 241 of confidence scores  $C_{qt}$ , since for any match in  $M_{qt}$  we only know the nearest  
 242 neighbours amongst the features of  $R_q$  and  $R_t$ . To get around this problem we  
 243 either assume to know one of the images beforehand (the *retrieval variation*) or  
 244 – in case we cannot make that assumption – we compute the features of one of  
 245 the images (the *general variation*). For a given match between features  $f_q$  and  
 246  $f_t$  we can now find the second closest neighbour  $f_b$  and calculate the  
 247 confidence as  $r = d(f_q, f_t)/d(f_q, f_b)$ .

248 In the retrieval scenario where we know an image beforehand, we can further  
 249 optimize this step. If we assume the *target image* is known in advance we can  
 250 approximate the ratio by letting  $\hat{r} = d(f_t, f_q)/d(f_t, f_b)$ . Here the feature  $f_b$  is  
 251 also part of the target image, which means we can pre-calculate  $d(f_t, f_b)$  for  
 252 all features in the cached *target image* before we start matching. When these  
 253 distances are tied to their corresponding feature in the target image, it becomes  
 254 trivial to calculate  $\hat{r}$ .

### 255 3.4 Exploring for Matches

256 In each iteration we compute a new set of seed matches, i.e. positions that  
 257 might yield more matches in the image. For each region  $R_i$  evaluated during the  
 258 collection step we now have a set of matches  $M_i$  and a set of confidence levels  $C_i$   
 259 that we can use to predict whether the neighbourhood of  $R_i$  is worth exploring.

260 There are many possible heuristics for predicting possible regions including  
 261 local and global epipolar assumptions and partial graph isomorphisms. However,  
 262 for the sake of simplicity and speed we have chosen a straightforward approach  
 263 based on weak angular assumptions, as illustrated in Figure 2. Assuming that the  
 264 *target image* has been cached either in advance or before the matching step, the  
 265 blue shaded areas show  $R_q$  and  $R_t$  for a given seed match. In the *target image* we  
 266 collect all features in a given radius while we compute features in the rectangular  
 267  $R_q$  in the *query image*. For performance reasons we compute all features in the  
 268

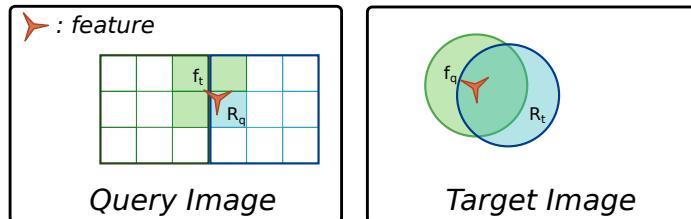


Fig. 2: Exploration of features based on a match. The areas shaded blue are the areas that were searched to obtain the match. The areas shaded green are candidate areas for obtaining more matches.

blue square but match only the features inside the shaded area. Next a match is found in the collection step between  $f_q$  and  $f_t$ . Based on the position of  $f_q$  in  $R_q$  we select three areas with potential for more matches. The center of each is matched with the center position  $f_t$  to produce three seed matches for the next iteration. In Figure 1 we illustrate this process in the *Fast-Match* image pair. The *query image* to the right in the figure has a small blue square for each region that has been used while matching the two images.

In practice it is necessary that these squares overlap in order to detect features lying close to the edges. This incurs a bit of overhead which is minimized by only collecting features for groups of 9 squares. In order to avoid computing the same matches or features twice, quite a bit of care has to be expended making sure that results are properly cached. A matrix containing ‘bins’ of features can conveniently be used to store features from different image regions. Similarly a hash-set is suitable for keeping track of which regions have already been matched and which matches have already been found.

### 3.5 Computational Complexity

The difference between the *general* and *retrieval* variation of *Fast-Match* consists in whether we compute the features in the *target image* (*general*) or if we presume the features are computed offline (*retrieval*). Computing the features is linear in the number of feature points, while finding  $d(f_q, f_b)$  for all  $n$  features in the *target image* can be done in  $O(n \log n)$  using metric trees.

Once the features and their distances have been computed, the algorithm iteratively finds new seed matches based on previous sets of seed matches. For each seed match we collect new matches, calculate confidence scores, and obtain new seed matches. Each of these steps varies only with the local region size which is constant. As a consequence the running time is linear in the amount of possible seed matches.

For the overall running time to be linear, we will need to show that the number of possible seed matches is linear in the number of image features. We will prove that this is the case on average and provide an upper bound for the probability that the average case will not happen. We assume that outside of true

315 correspondences, features have better or at least equally good matches in terms  
 316 of confidence in the image they come from. More rigorously put: For any feature  
 317 in the target image  $f_t$  let  $f_1, f_2, \dots$  be the best, second best, etc matching feature  
 318 from either image which is not a true correspondence. Let  $A_{ti}$  be the event that  
 319  $f_i$  is found in the *query image*, then  $\mathbb{P}(A_{ti}) \leq 0.5$  and  $A_{ti}$  is independent from  
 320  $A_{tj}$  when  $i \neq j$ .

321 For a given feature in the *target image*  $f_t$ , its nearest neighbour in the *target*  
 322 *image*  $f_b$ , and the set of features in the *query image*  $F_{query}$ , we can define the  
 323 stochastic variable  $X_t$  as follows:

$$324 \quad X_t = |\{f_q \mid f_q \in F_{query}, d(f_q, f_t) < d(f_q, f_b)\}| \quad (1)$$

325 That is,  $X_t$  is the number of features in the query image closer to  $f_q$  than  $f_b$ .  
 326 For each feature  $f_t$ ,  $X_t$  is an i.i.d. stochastic variable.

327 In the worst case, each feature in the *target image* has a true correspondence  
 328 in the *query image* and  $\mathbb{P}(A_{ti} = 0.5)$ , in which case we compute the probability  
 329 that  $X_t = n$  as well as the expected value and variance of  $X_t$ :

$$330 \quad \mathbb{P}(X_t = n) = \mathbb{P}(A_{t1} \wedge \dots \wedge A_{tn}) = 0.5^n \quad (2)$$

$$331 \quad \mathbb{E}(X_t) = \sum_{i=1}^{\infty} i0.5^i = 2 \quad (3)$$

$$332 \quad \text{Var}(X_t) = \sum_{i=1}^{\infty} 0.5^i(i-2)^2 = 6 \quad (4)$$

333 For  $\tau = 1$  we accept a seed match when  $d(f_t, f_q) \leq d(f_t, f_b)$ . That means  
 334 that for any feature in the *query image*, we accept at most  $X_t$  matches. To find  
 335 the total amount of seed matches given  $n$ , we let  $S_n = X_1 + \dots + X_n$ , in which  
 336 case  $\mathbb{E}(S_n) = 2n$ . This proves that the expected amount of seed matches is linear  
 337 in the number of target features.

338 For the above case it is theoretically possible that for every one of the  $n$   
 339 features we end up considering  $kn$  seed matches for some constant  $k$ , i.e. a  
 340 quadratic behaviour. Using Chebyshev's inequality we can provide an upper  
 341 bound on the probability of this event:

$$342 \quad \mathbb{P}\left(\left|\frac{S_n}{n} - \mathbb{E}(X_t)\right| \geq kn\right) \leq \frac{\text{Var}\left(\frac{S_n}{n}\right)}{k^2 n^2} = \frac{6}{k^2 n^3} \quad (5)$$

343 It is clear that for any constant  $k$  we can pick a number of features  $n$  which  
 344 render the likelihood of a quadratic behaviour infinitesimal. In essence the larger  
 345  $n$  becomes, the less likely *Fast-Match* is to exhibit super linear behaviour.

346 The noteworthy part of this result is that for the *retrieval* variation we can  
 347 match two images in  $O(n)$ . Since there are no large constants involved this makes  
 348 it possible to rapidly match very large images. In particular *Fast-Match* is suited  
 349 for cases where we are looking to match several large *query images* to one *target*  
 350 *image*. In this case we can compute the features and distances of the *target*

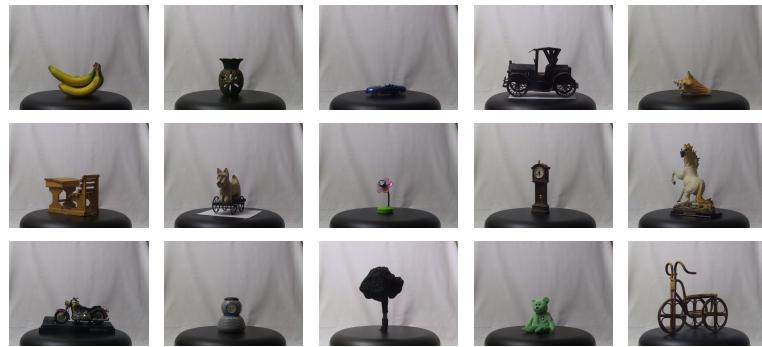
360     *image* and then match each *query image* in linear time. In contrast the general  
361     variation still has a complexity of  $O(n \log n)$  due to the necessity of calculating  
362     distances online.

## 364     4 Experimental Setup

366     We evaluate *Fast-Match* over 3024 image pairs featuring various 3D objects at  
367     different angles to measure precision and recall, as well as two pairs of images  
368     with high pixel counts to measure performance in terms of speed. We compare  
369     the algorithm to the standard *Ratio-Match* [1] as well as the newer *Mirror-*  
370     *Match* [18]. These two algorithms were selected because they are magnitudes  
371     faster than the fastest geometric algorithms while at the same time providing  
372     robust matches.

### 374     4.1 Evaluation of Fast-Match on 3D Objects

376     The 3D Objects dataset by Moreels and Perona [3] allows us to experimentally  
377     compare matching algorithms over a large range of object and surface types  
378     rotated on a turnstile and photographed from every angle in increments of 5  
379     degrees. 15 objects from the dataset are shown in Figure 3. We use images of  
380     84 different objects photographed under three different lighting conditions at  
381     12 different angle intervals, conducting experiments with a total of 3024 image  
382     pairs. All photos a taken with a consumer camera in 3.1 mega pixel resolution.



396     Fig. 3: 15 objects from the 3D Objects dataset by Moreels and Perona [3].

399     To validate matches, Moreels and Perona propose a method using epipolar  
400     constraints [3, p.266]. According to their experiments, these constraints are able  
401     to identify true correspondences with an error rate of 2%. We use their proposed  
402     method to generate the ground truth for the evaluation of our framework.

403     To compute the total number of possible true correspondences, we take each  
404     feature in a *query image* and count how many of them have a feature in the *target*

405 *image* which would satisfy the epipolar constraints. When using this dataset,  
406 features with no correspondences were not included in the set of features for  
407 testing, so as to avoid matching non-moving background and foreground objects.

408 We evaluate all matching algorithms from our framework on the 3D Objects  
409 dataset by matching images at different angular intervals. For each object we  
410 pick the *query* image as the image taken at 10 degrees rotation for calibration  
411 stability. We then match this image with the same object turned an additional  $\Delta$   
412 degrees,  $\Delta \in \{5, 10, \dots, 60\}$ . For every angle interval we compare images taken  
413 under 3 different lighting conditions as provided by the dataset. We include all  
414 objects in the database for which photos at 5 degree angle intervals are available  
415 except for the “Rooster” and “Sponge” objects due to image irregularities.

416

#### 417 4.2 Configuration of *Fast-Match*

418

419 The central parameters of *Fast-Match* are the confidence threshold  $\tau$  for select-  
420 ing seed matches and the final confidence threshold applied to the total set of  
421 matches. For the experiments on the 3D Objects we let  $\tau = 0.9$  and created  
422 precision/recall plots by varying the confidence threshold over the final set of  
423 matches.

424 To achieve a good balance between speed and robustness we let the region in  
425 which we extract features be a square with a side length of 90 pixels. On empirical  
426 tests we found that anything smaller would result in decreased performance while  
427 much bigger regions would decrease speed. The 90 pixel window is split in to  
428 nine smaller squares (see also Figure 2). When a seed match falls in any of these  
429 squares we match only the features within the smaller square. We let both the  
430 regions and the smaller squares overlap each other at all edges with 25 pixels in  
431 order to capture feature points lying close to an edge. For the image we have  
432 already cached we find all features within a radius of 50 pixels of the seed match.

433

## 434 5 Results

435

436 Figure 4 shows the performance of the different matching methods in our pro-  
437 posed framework for 12 increasingly bigger angle differences. The results are  
438 shown in a precision-recall plot to make it easy to compare performance in terms  
439 of precision at similar levels of recall. For each plot we show the accumulated re-  
440 sults on all 3D objects, weighted by the number of possible true correspondences  
441 for the individual objects. This ensures that each object contributes equally to  
442 the final result, despite some objects resulting in disproportionately more matches  
443 than others.

444 At low angular differences all algorithms show similar performance at low re-  
445 call, however *Fast-Match* is superior to *Ratio-Match* and *Mirror-Match* at higher  
446 recall, more than doubling the precision at similar recall levels. This picture re-  
447 mains the same at higher angular differences, but with a higher performance  
448 gap between *Fast-Match* and the other algorithms at low recall while more than  
449 doubling precision at higher recalls. For image pairs with angular differences of

405

406

407

408

409

410

411

412

413

414

415

416

417

418

419

420

421

422

423

424

425

426

427

428

429

430

431

432

433

434

435

436

437

438

439

440

441

442

443

444

445

446

447

448

449

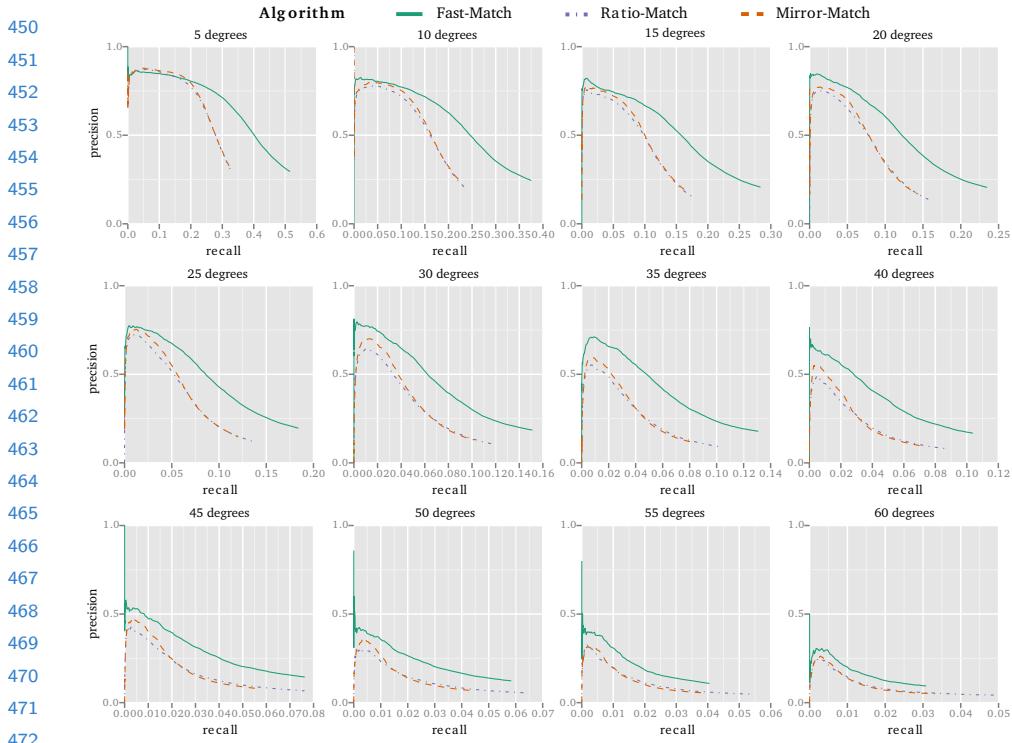


Fig. 4: Results for the 3D Objects dataset. Each plot contains data accumulated from 84 objects photographed under 3 different lighting conditions.

more than  $40^\circ$  the performance of all algorithms declines, with *Fast-Match* still coming out on top.

The high recall rates of *Fast-Match* are partially explained by the fact that we can allow our confidence thresholds to be more lenient when we already know that we are likely to find true correspondences in the regions that we are matching. However, lenient thresholds can easily impact the matching speed, so in practical situations we usually have to choose between matching fast and precisely, or more slowly with higher recall.

Figure 5 demonstrate the speed of *Fast-Match* on image sizes as is typically produced by modern consumer cameras and camera phones. In Figure 5 we compare the speed of the two variations of *Fast-Match* to different variations of *Ratio-Match* on two image pairs, one with two images of 10 megapixel and another where the same image pair has been scaled to 3 megapixels. For the larger images we see the clearest difference in performance. The retrieval variant of *Fast-Match* matches the image pairs in around one second while a retrieval variant of *Ratio-Match* with precomputed features spends eight seconds on the same task. This result is roughly equal to the time the general variation of *Fast-Match* spends matching the two images without any pre-computations. For the nearest

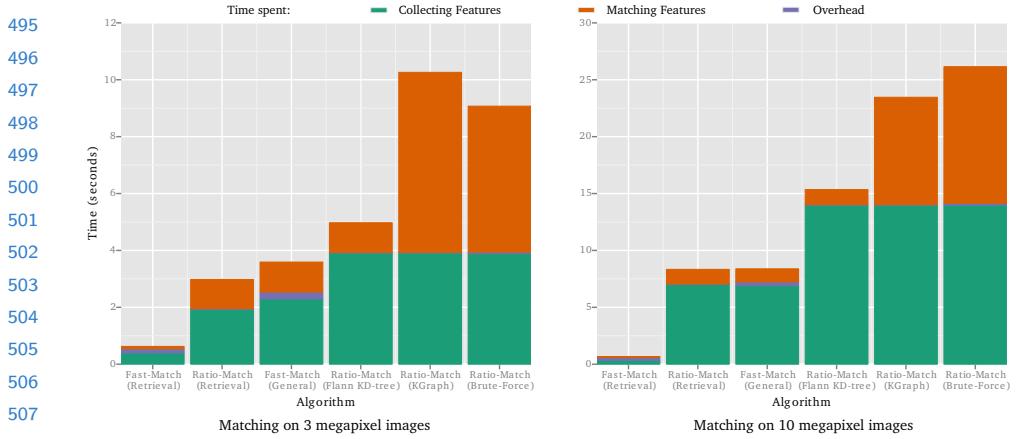


Fig. 5: Running times of different variations of *Fast-Match* and *Ratio-Match*. The left plot shows results for two 3 Mpixel images while the plot on the right shows results for two 10 Mpixel images. The two leftmost results in either plot presume a precached image while the rest require no caching. The image pairs used can be seen in Figure 1.

neighbor search we note that for large images there are substantial speed gains to be had by choosing the right algorithm for finding nearest neighbors, with Muja et al.'s Flann matcher being vastly superior to brute force and KGraph.

For the 3 megapixel picture the pattern repeats itself, although with smaller margins separating *Fast-Match* and *Ratio-Match*. When images are smaller, computations like obtaining the initial seed matches are comparatively more costly.

Figure 1 shows the result from matching the pair of three megapixel images using *Fast-Match* and *Ratio-Match*. *Fast-Match* ends up only matching the part of the images that are likely to have correspondences. This highlights both a strength and a weakness of *Fast-Match*. It is this myopic matching approach that is responsible for *Fast-Match*' speed and precision; on the other hand, *Fast-Match* is not the best candidate for images that are almost identical, because its step by step approach incurs too much overhead to compete in speed with *Ratio-Match* in those cases.

## 6 Conclusion

We have introduced *Fast-Match* in two forms, a general and a retrieval variation. The retrieval variation assumes that we know one of the images that we are matching ahead of time in order to pre-cache feature points. The general variation has no assumptions and performs the pre-caching of feature points online instead. We have proven that while the general variation has a computational complexity of  $O(n \log n)$  with  $n$  being the number of feature points, the retrieval variation will on average run in  $O(n)$ . From experiments on large images we

540 have shown that *Fast-Match* can be a magnitude faster than *Ratio-Match* while  
541 providing comparable results. We further compared *Fast-Match* to *Ratio-Match*  
542 and *Mirror-Match* using images of rotated 3D objects to demonstrate that *Fast-  
543 Match* outperforms the other algorithms significantly over 3024 image pairs and  
544 in most cases doubles the matching precision at similar recall rates. *Fast-Match*  
545 is particularly applicable in cases where we are interested in matching one image  
546 to multiple large images, in which case we can quickly pre-cache the image and  
547 use this data for each of the other images.

548

549

550 

## References

551

- 552 1. Lowe, D.G.: Distinctive image features from scale-invariant keypoints. International  
553 Journal on Computer Vision **60** (2004) 91–110
- 554 2. Mikolajczyk, K., Schmid, C.: A performance evaluation of local descriptors. IEEE  
555 Trans. Pattern Analysis and Machine Intelligence **27** (2005) 1615–1630
- 556 3. Moreels, P., Perona, P.: Evaluation of features detectors and descriptors based on  
557 3D objects. International Journal on Computer Vision **73** (2007) 263–284
- 558 4. Heinly, J., Dunn, E., Frahm, J.M.: Comparative evaluation of binary features. In:  
559 Proc. European Conference on Computer Vision (ECCV). (2012) 759–773
- 560 5. Bay, H., Tuytelaars, T., Van Gool, L.: SURF: Speeded up robust features. In:  
561 Proc. European Conference on Computer Vision (ECCV). (2006) 404–417
- 562 6. Calonder, M., Lepetit, V., Strecha, C., Fua, P.: BRIEF: binary robust independent  
563 elementary features. In: Proc. European Conference on Computer Vision (ECCV).  
564 (2010) 778–792
- 565 7. Leutenegger, S., Chli, M., Siegwart, R.Y.: BRISK: Binary robust invariant scalable  
566 keypoints. In: Proc. International Conference on Computer Vision (ICCV), IEEE  
567 (2011) 2548–2555
- 568 8. Ke, Y., Sukthankar, R.: PCA-SIFT: A more distinctive representation for local  
569 image descriptors. In: Proc. Conference on Computer Vision and Pattern Recog-  
570 nition (CVPR). Volume 2., IEEE (2004) II–506
- 571 9. Li, J., Allinson, N.M.: A comprehensive review of current local features for com-  
572 puter vision. Neurocomputing **71** (2008) 1771–1787
- 573 10. Beis, J.S., Lowe, D.G.: Shape indexing using approximate nearest-neighbour search  
574 in high-dimensional spaces. In: Proc. Conference on Computer Vision and Pattern  
575 Recognition (CVPR), IEEE (1997) 1000–1006
- 576 11. Lowe, D.G.: Object recognition from local scale-invariant features. In: Proc. Inter-  
577 national Conference on Computer Vision (ICCV). Volume 2., Ieee (1999) 1150–1157
- 578 12. Muja, M., Lowe, D.G.: Fast approximate nearest neighbors with automatic algo-  
579 rithm configuration. In: Proc. International Joint Conference on Computer Vision,  
580 Imaging and Computer Graphics Theory and Applications (VISAPP). (2009) 331–  
581 340
- 582 13. Dong, W., Moses, C., Li, K.: Efficient k-nearest neighbor graph construction for  
583 generic similarity measures. In: Proc. International Conference on World Wide  
584 Web (IW3C2), ACM (2011) 577–586
- 585 14. Deriche, R., Zhang, Z., Luong, Q.T., Faugeras, O.: Robust recovery of the epipo-  
586 lar geometry for an uncalibrated stereo rig. In: Proc. European Conference on  
587 Computer Vision (ECCV). (1994) 567–576

14 ACCV-14 submission ID 325

- 585 15. Baumberg, A.: Reliable feature matching across widely separated views. In: Proc.  
586 Conference on Computer Vision and Pattern Recognition (CVPR). Volume 1.  
587 (2000) 774–781  
588 16. Rabin, J., Delon, J., Gousseau, Y.: A statistical approach to the matching of local  
589 features. SIAM Journal on Imaging Sciences **2** (2009) 931–958  
590 17. Brown, M., Szeliski, R., Winder, S.: Multi-image matching using multi-scale ori-  
591 ented patches. In: Proc. Conference on Computer Vision and Pattern Recognition  
(CVPR). Volume 1. (2005) 510–517  
592 18. Arnfred, J.T., Winkler, S., Süsstrunk, S.: Mirror Match: Reliable feature point  
593 matching without geometric constraints. In: Proc. IAPR Asian Conference on  
594 Pattern Recognition (ACPR). (2013) 256–260  
595 19. Barnes, C., Shechtman, E., Finkelstein, A., Goldman, D.: PatchMatch: a random-  
596 ized correspondence algorithm for structural image editing. ACM Transactions on  
597 Graphics-TOG **28** (2009) 24  
598 20. Fischler, M.A., Bolles, R.C.: Random sample consensus: a paradigm for model  
599 fitting with applications to image analysis and automated cartography. Commu-  
600 nications of the ACM **24** (1981) 381–395  
601 21. Torr, P.H.S., Zisserman, A.: MLESAC: A new robust estimator with application  
602 to estimating image geometry. Computer Vision and Image Understanding **78**  
603 (2000) 138–156  
604 22. Szeliski, R.: Computer Vision: Algorithms and Applications. Springer (2010)  
605 23. Kim, J., Choi, O., Kweon, I.S.: Efficient feature tracking for scene recognition  
606 using angular and scale constraints. In: Proc. IEEE/RSJ International Conference  
607 on Intelligent Robots and Systems (IROS), Nice, France (2008) 4086–4091  
608 24. Schmid, C., Mohr, R.: Local grayvalue invariants for image retrieval. IEEE Trans-  
609 actions on Pattern Analysis and Machine Intelligence **19** (1997) 530–535  
610 25. Torresani, L., Kolmogorov, V., Rother, C.: Feature correspondence via graph  
611 matching: Models and global optimization. In: Proc. European Conference on  
612 Computer Vision (ECCV). Springer (2008) 596–609  
613 26. Yarkony, J., Fowlkes, C., Ihler, A.: Covering trees and lower-bounds on quadratic  
614 assignment. In: Proc. Conference on Computer Vision and Pattern Recognition  
(CVPR), IEEE (2010) 887–894  
615 27. Cho, M., Lee, J., Lee, K.M.: Reweighted random walks for graph matching. In:  
616 Proc. European Conference on Computer Vision (ECCV). Springer (2010) 492–505  
617 28. Leordeanu, M., Hebert, M.: A spectral technique for correspondence problems  
618 using pairwise constraints. In: Proc. Conference on Computer Vision and Pattern  
619 Recognition (CVPR). Volume 2., San Diego, CA (2005) 1482–1489  
620 29. Yuan, Y., Pang, Y., Wang, K., Shang, M.: Efficient image matching using weighted  
621 voting. Pattern Recognition Letters **33** (2012) 471–475  
622 30. Pang, Y., Shang, M., Yuan, Y., Pan, J.: Scale invariant image matching using  
623 triplewise constraint and weighted voting. Neurocomputing **83** (2012) 64–71  
624 31. Cho, M., Lee, J., Lee, K.M.: Feature correspondence and deformable object match-  
625 ing via agglomerative correspondence clustering. In: Proc. International Conference  
626 on Computer Vision (ICCV), IEEE (2009) 1280–1287  
627 32. Wu, L., Niu, Y., Zhang, H., Zhu, H., Shen, L.: Robust feature point matching  
628 based on local feature groups (LFGs) and relative spatial configuration. Journal  
629 of Computational Information Systems **7** (2011) 3235–3244  
629 33. Chen, H.Y., Lin, Y.Y., Chen, B.Y.: Robust feature matching with alternate hough  
and inverted hough transforms. In: Proc. Conference on Computer Vision and  
Pattern Recognition (CVPR), IEEE (2013) 2762–2769